

# 0-to-hero

07/10 Mentors <> 07/10 Sessions

# Elkhan Shahverdi

— — —

Senior Software Engineer at Insparx Gmbh

[https://www.linkedin.com/in/elkhan-shahverdi-59356166/  
shahverdiyev90@gmail.com](https://www.linkedin.com/in/elkhan-shahverdi-59356166/shahverdiyev90@gmail.com)

Twitter: @shahverdiyev90

Ankara University – Computer Engineering (BSc)

University of Tartu – Software Engineering (MSc)

# Estonia

— — —

Study:

<https://www.studyinestonia.ee/en>

Work:

<https://www.workinestonia.com/>

<https://devhunt.ee>

# Stream Processing

— — —

## Agenda

- Derived data
- Stream processing
- Apache Flink Architecture
- Benchmark of Stream Processing engines.


# Derived Data

— — —

- Network Monitoring and Traffic engineering
- Sensor logs
- Telecom call-detail records
- Mobile Application logs
- Social network trends
- Website logs and clickstreams
- Other massive data sets...

# Importance of real-time processing


— — —



# What is Stream Processing

— — —


- Processing data in motion
- Querying continuous data stream
- Detect conditions fast within a small time period
- Related synonyms for Stream Processing are: Real-time analytics, Streaming analytics, Complex Event Processing, Event Processing



# What is Stream Processing doing?

— — —


- Sourcing
- Filtering
- Transformation
- Map-Reduce
- Grouping
- Windowing
- Stream-Stream Join
- Stream-Table Join
- Sink






# Windowing


Fixed / Tumbling Window



Sliding / Hopping Window



Session Window



# API

— — —

- Programmatic

```
messageStream
    .rebalance()
    .flatMap(new DeserializeBolt())
    .filter(new EventFilterBolt())
    .<Tuple2<String, String>>project(2, 5)
    .flatMap(new RedisJoinBolt())
    .keyBy(0)
    .flatMap(new CampaignProcessor());
```

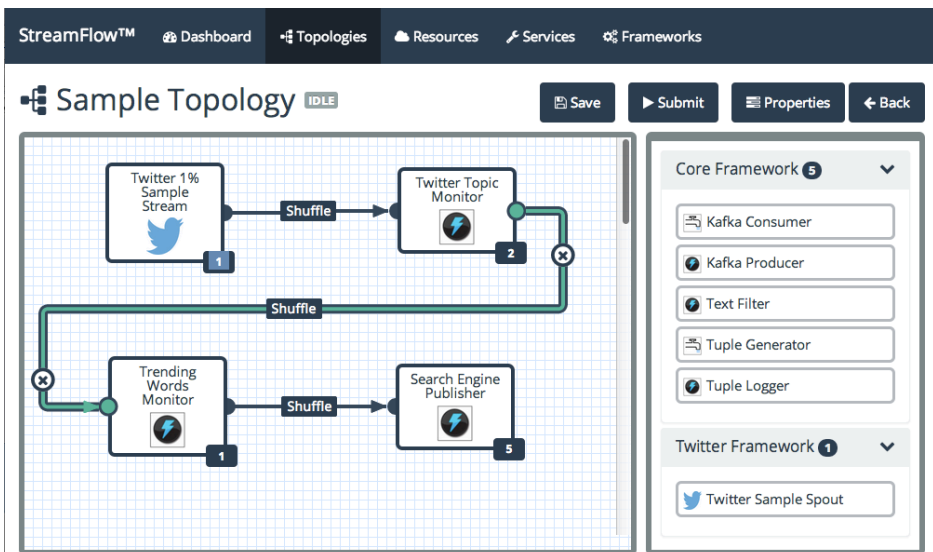
- Streaming SQL

```
SELECT STREAM CEIL(rowtime TO HOUR) AS rowtime,
    productId,
    COUNT(*) AS c,
    SUM(units) AS units
FROM Orders
GROUP BY CEIL(rowtime TO HOUR), productId;
```

# API

- GUI-based / drag and drop

- Declarative (Json, Yaml)




```
#STORM Specific
storm.workers: 72
storm.ackers: 9
storm.ack: "enabled"

#SPARK Specific
spark.batchtime: 3000
spark.master: "spark://stream-node01:7077"
spark.app.name: "KafkaRedisAdvertisingStream"
```

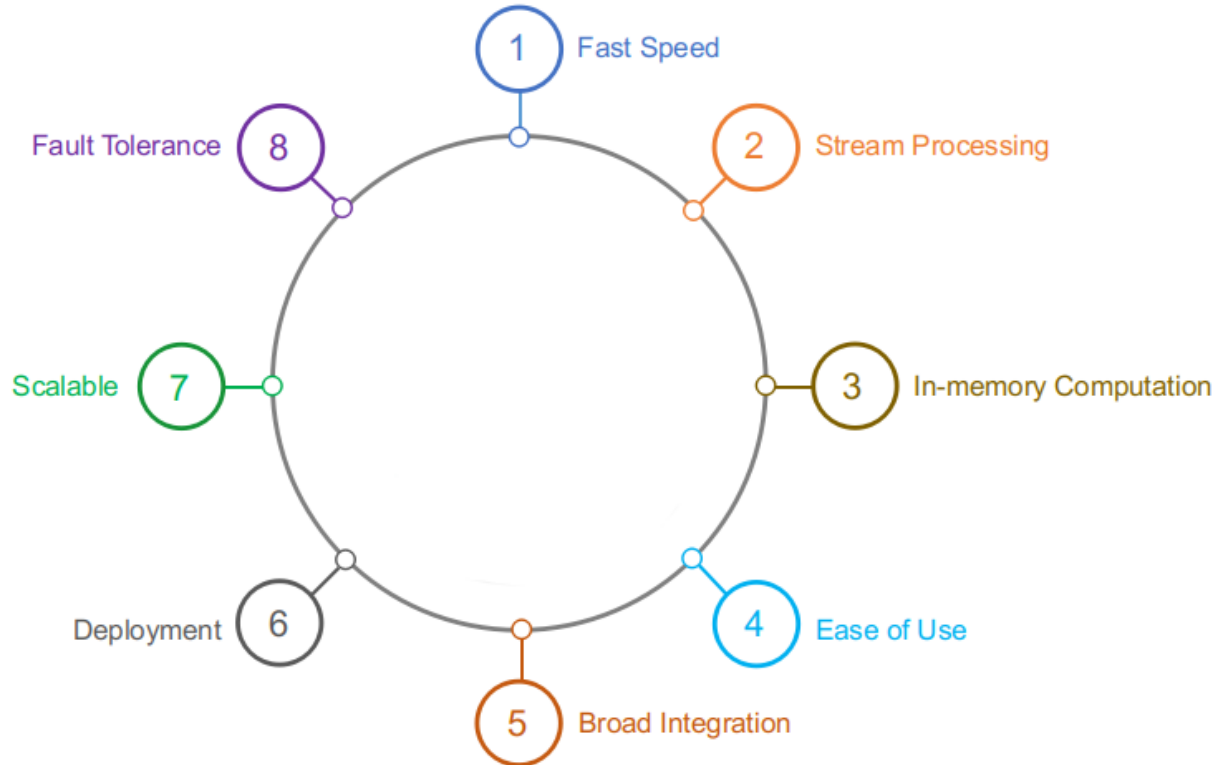
# Stream Processing Engines

---



# Features of Stream Processing Engines

— — —



# Delivery Guarantees

— — —

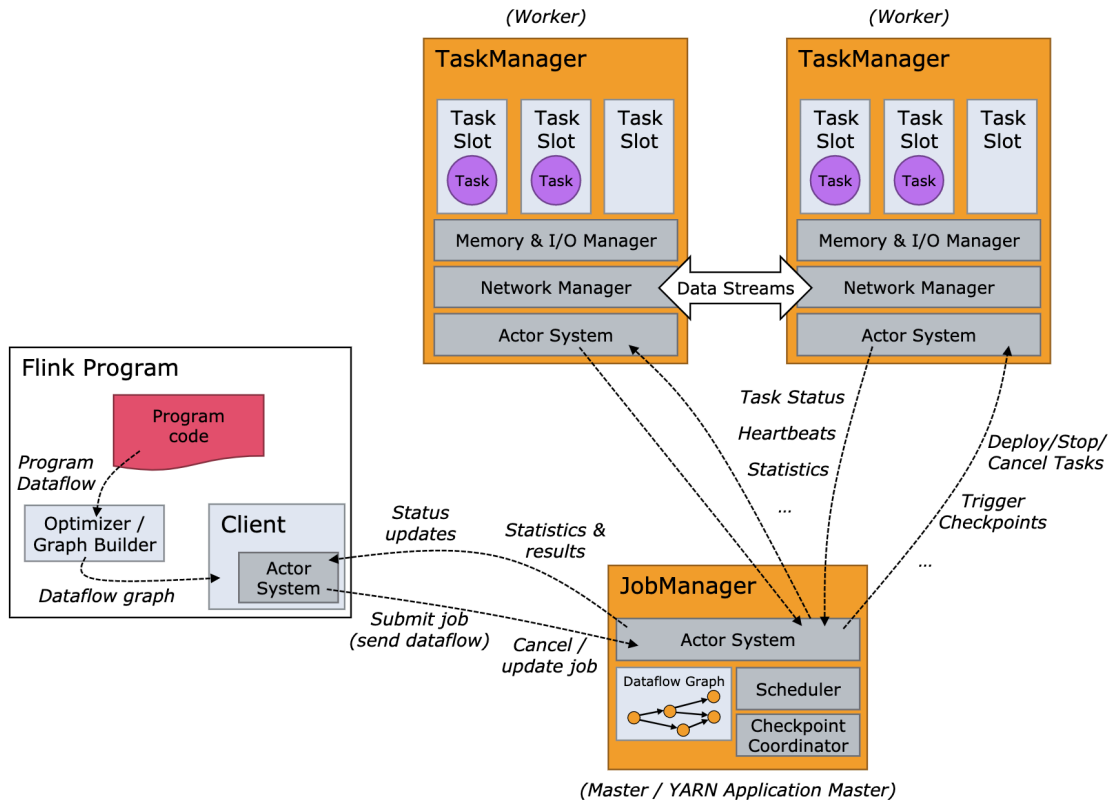
- At most once - fire and forget
- At least once
- Exactly once

# Apache Flink Architecture

---

- Program
  - It is a piece of code, which you run on the **Flink Cluster**.
- Client
  - It is responsible for taking code (program) and constructing job dataflow graph, then passing it to **JobManager**.
- Job Manager
  - Also called **masters**.
- Task Manager
  - Also called **workers** or **slaves**.

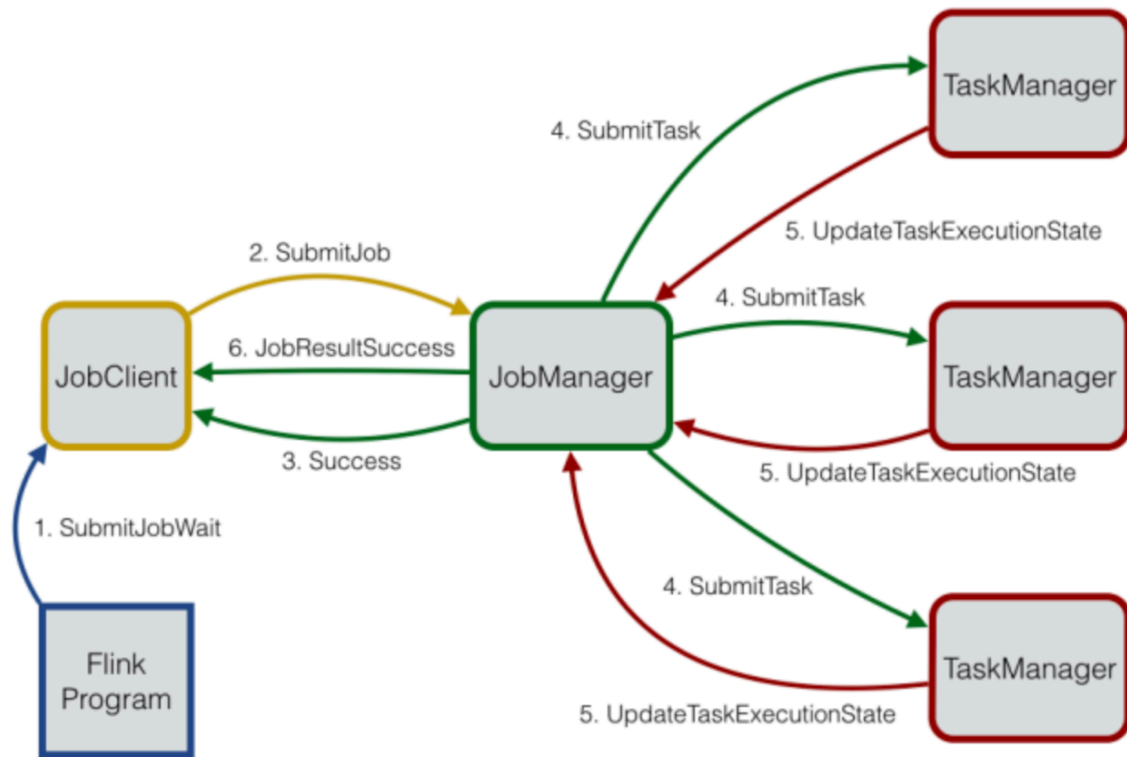
# Apache Flink Architecture



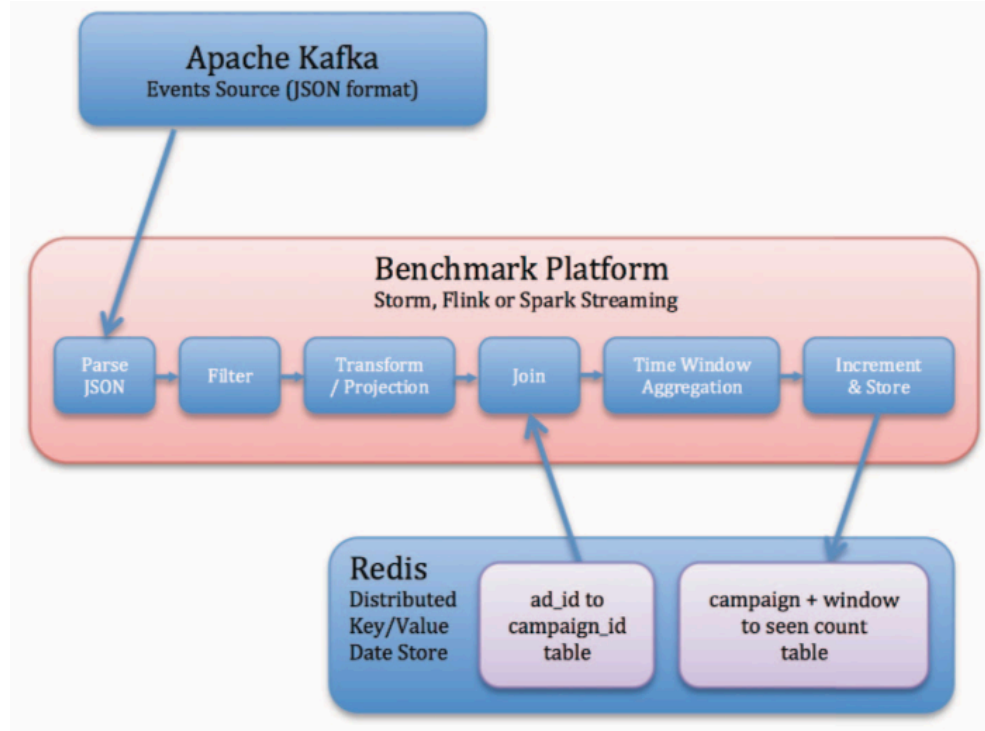


# Apache Flink Architecture

```
messageStream
  .rebalance()
  .flatMap(new DeserializeBolt())
  .filter(new EventFilterBolt())
  .<Tuple2<String, String>>project(2, 5)
  .flatMap(new RedisJoinBolt())
  .keyBy(0)
  .flatMap(new CampaignProcessor());
```



# Experiment Design



1. C. Sanket, D. Derek, E. Bobby, F. Reza, G. Thomas, H. Mark, L. Zhuo, N. Kyle, P. Kishorkumar, J. P. Boyang ja P. Paul, „Benchmarking Streaming Computation Engines: Storm, Flink and Spark Streaming.,“ IEEE, 2016.

# Comparison with Yahoo Streaming Benchmark

— — —


Tools	Yahoo's version (2015)	My version (2018)
Flink	1.1.3	1.5.0
Spark DStream	1.6.2	2.3.0
Strom	0.9.7	1.2.1
Redis	3.0.5	4.0.8
Kafka Broker	0.8.2.1	0.11.0.2
Spark Structured Streaming	-	2.3.0
Kafka Stream	-	1.1.0

Benchmark	Benchmark Metrics	Servers Type	Servers CPU	Servers memory
Yahoo's version	Latency, Throughput	Dedicated server	16 cores (8 physical, 16 hyperthreading)	24GB memory
My Version	Latency, Throughput, Resource Consumption	Virtual Private Server	16 virtual cores	32GB memory

GitHub Repository of the project <https://github.com/elkhan-shahverdi/streaming-benchmarks>


# Experiment Topology

- 10 Stream Processing nodes
- 5 Kafka nodes
- 3 Zookeeper nodes
- 10 Producer nodes
- 1 Redis node




— — —


# Experiment Results




(a) Storm




(b) Flink




(c) Spark DStream



(d) Spark Structured Streaming




(e) Kafka Streams




(f) Hazelcast Jet


Fig. 2. Latency versus percentage of completed tuples



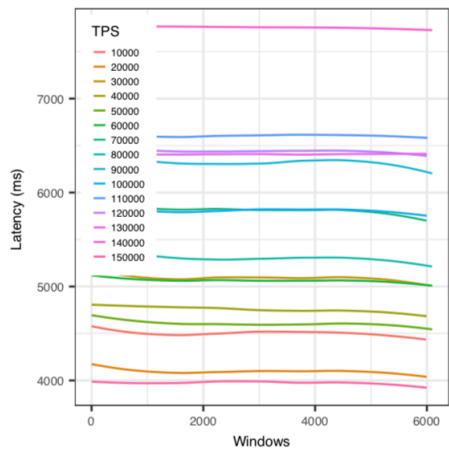
(a) Storm




(b) Flink




(c) Spark DStream



(d) Spark Structured Streaming




(e) Kafka Streams




(f) Hazelcast Jet

Fig. 3. Loess Regression of Latencies




(a) 90% Percentile Latency




(b) 99% Percentile Latency




— — —



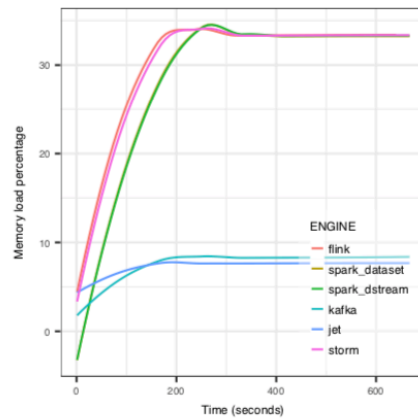
(a) Streaming servers CPU load



(b) Kafka servers CPU load




(c) Streaming servers memory consumption




(d) Kafka servers memory consumption

Fig. 5. Resource Consumption at 60K TPS


— — —




(a) Streaming servers CPU load



(b) Kafka servers CPU load



(c) Streaming servers memory consumption



(d) Kafka servers memory consumption

Fig. 6. Resource Consumption at 150K TPS

# Experiment Conclusion

— — —

- Flink and Kafka are most noticeable real-time stream processors.
- Flink has lower latency than Kafka.
- Apache Storm latency can race with Flink with the low amount of data.
- Flink custom memory management is not better than Java GC.
- Spark Structured Streaming is the more resistant than DStream under huge amount of data.
- Flink owes the highest performance for more resource consumptions

# Conclusion

---

- Principle is important than tools

# References

— — —

1. [Big Stream Processing Systems: An Experimental Evaluation](#)  
by Elkhan Shahverdi, Ahmed Awad, Sherif Sakr
2. [Designing Data-Intensive Applications.](#) By Kleppmann,  
Martin
3. [Streaming Systems: The What, Where, When, and How of  
Large-Scale Data Processing.](#) by Tyler Akidau, Slava  
Chernyak, Reuven Lax
4. <https://cwiki.apache.org/confluence/display/FLINK/Flink+Internals>
5. <https://github.com/elkhan-shahverdi/streaming-benchmarks>